

SmartInternz Externships

Applied Data Science

Assignment- 3

Name : MITHUN JOGI R

Reg no : 20BCE7163

Campus: VIT-AP

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

2) Load the dataset

```
In [2]: df = pd.read_csv("housing.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditic
0	13300000	7420	4	2	3	yes	no	no		no
1	12250000	8960	4	4	4	yes	no	no		no
2	12250000	9960	3	2	2	yes	no	yes		no
3	12215000	7500	4	2	2	yes	no	yes		no
4	11410000	7420	4	1	2	yes	yes	yes		no

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null    int64
1   area                 545 non-null    int64
2   bedrooms             545 non-null    int64
3   bathrooms             545 non-null    int64
4   stories              545 non-null    int64
5   mainroad             545 non-null    object
6   guestroom            545 non-null    object
7   basement             545 non-null    object
8   hotwaterheating      545 non-null    object
9   airconditioning      545 non-null    object
10  parking              545 non-null    int64
11  furnishingstatus     545 non-null    object
dtypes: int64(6), object(6)
```


3. Perform Below Visualizations.

- Univariate Analysis
- Bi - Variate Analysis
- Multi - Variate Analysis

univariate analysis

```
In [5]: sns.histplot(df['area'])
```


```
Out[5]: <Axes: xlabel='area', ylabel='Count'>
```



```
sns.histplot(df['price'])
```

```
In [6]: <Axes: xlabel='price', ylabel='Count'>
```

```
Out[6]:
```




```
In [7]: sns.countplot(x = df['bedrooms'])
```

```
Out[7]: <Axes: xlabel='bedrooms', ylabel='count'>
```


```
In [8]:
```

```
Out[8]:
```

```
<Axes:
xlabel=
'bathrooms',
ylabel=
'count'
>
```



```
sns.countplot(x = df['bathrooms'])
```




```
In [9]: sns.countplot(x = df['stories'])
```

```
Out[9]: <Axes: xlabel='stories', ylabel='count'>
```


```
In [10]:
```

```
Out[10]:
```

```
<Axes:
xlabel=
'guestr
oom',
ylabel=
'count'
>
```



```
sns.countplot(x = df['guestroom'])
```




```
In [11]: sns.countplot(x = df['mainroad'])
```

```
Out[11]: <Axes: xlabel='mainroad', ylabel='count'>
```


```
In [12]:
```

```
Out[12]:
```

```
<Axes:
xlabel=
'parkin
g',
ylabel=
'count'
>
```




```
sns.countplot(x = df['parking'])
```

```
In [13]: sns.countplot(x = df['furnishingstatus'])
```

```
Out[13]: <Axes: xlabel='furnishingstatus', ylabel='count'>
```




bivariate analysis

```
In [14]:
```

```
Out[14]:
```

```
sns.scatterplot(data = df, x = 'area', y = 'price')
```

```
<Axes:  
xlabel=  
'area',  
ylabel=  
'price'  
>
```



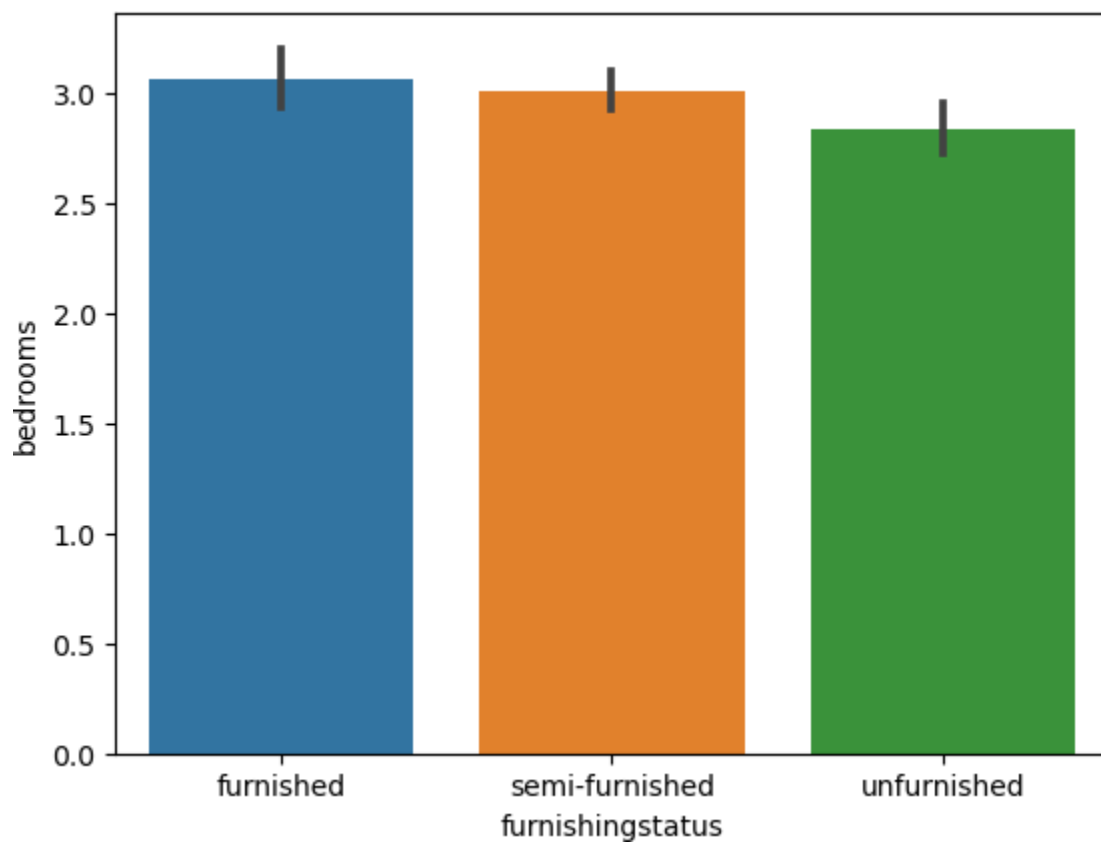
```
In [15]: sns.barplot(data = df, x = 'furnishingstatus', y = 'bedrooms')
```

```
Out[15]: <Axes: xlabel='furnishingstatus', ylabel='bedrooms'>
```


```
In [16]:
```

```
Out[16]:
```

```
<Axes:
xlabel=
'furnis
hingsta
tus',
ylabel=
'area'>
```



```
sns.barplot(data = df, x = 'furnishingstatus', y = 'area')
```




```
In [17]: sns.barplot(data = df, x = 'furnishingstatus', y = 'price')
```

```
Out[17]: <Axes: xlabel='furnishingstatus', ylabel='price'>
```


```
In [18]:
```

```
Out[18]:
```

```
<Axes:
xlabel=
'furnis
hingsta
tus',
ylabel=
'count'
>
```




```
sns.countplot(x = df['furnishingstatus'], hue = df['airconditioning'])
```



```
In [19]: sns.countplot(x = df['furnishingstatus'], hue = df['hotwaterheating'])
```

```
Out[19]: <Axes: xlabel='furnishingstatus', ylabel='count'>
```




multivariate analysis

```
In [20]:
```

```
Out[20]:
```

```
sns.heatmap(df.corr(numeric_only=True), annot = True)
```

```
<Axes:  
>
```

4. Perform descriptive statistics on the dataset.

```
In [21]: df.describe()
```

```
Out[21]:
```

	price	area	bedrooms	bathrooms	stories	parking
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	0.693578
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	0.861586
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	0.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	0.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	0.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	1.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	3.000000

5. Handle the Missing values.

```
In [22]: df.isnull().sum()
```


```
Out[22]: price      0
         area      0
         bedrooms  0
         bathrooms 0
         stories   0
```

```
mainroad      0
guestroom     0
basement      0
hotwaterheating 0
airconditioning 0
parking       0
furnishingstatus 0
dtype: int64
```

6. Find the outliers and replace the outliers

```
In [23]: sns.boxplot(df['price'])
```

```
Out[23]: <Axes: >
```




```
sns.boxplot(df['area'])
```

```
<Axes: >
```


```
In [24]:
```

```
Out[24]:
```



```
In [25]: median_age = df['price'].median()
df["price"] = np.where(df["price"] > 12000000, median_age, df['price'])
sns.boxplot(df['price'])
```


Out[25]: <Axes: >



```
In [26]: median_area = df['area'].median()
df["area"] = np.where(df["area"] > 14000, median_area, df['area'])
sns.boxplot(df['area'])
```

Out[26]:

<Axes:
>



7. Check for Categorical columns and perform encoding.

```
In [27]: from sklearn.preprocessing import OneHotEncoder
```

```
In [28]: encoding = pd.get_dummies(df, columns = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'furnishingstatus'])
```

```
In [29]: encoding.head()
```

```
Out[29]:
```

	price	area	bedrooms	bathrooms	stories	parking	mainroad_no	mainroad_yes	guestroom_no	guestroom_yes
0	4340000.0	7420.0	4	2	3	2	0	1	1	
1	4340000.0	8960.0	4	4	4	3	0	1	1	
2	4340000.0	9960.0	3	2	2	2	0	1	1	
3	4340000.0	7500.0	4	2	2	3	0	1	1	
4	11410000.0	7420.0	4	1	2	2	0	1	0	

8. Split the data into dependent and independent variables

```
df.columns
```

```
In [30]: Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking', 'furnishingstatus'], dtype='object')
```

```
Out[30]:
```

```
In [65]: # independent variables
X = encoding.drop(['price'], axis = 1)
X.head()
```

```
Out[65]:
```

	area	bedrooms	bathrooms	stories	parking	mainroad_no	mainroad_yes	guestroom_no	guestroom_yes	total_rooms
0	7420.0	4	2	3	2	0	1	1	0	8
1	8960.0	4	4	4	3	0	1	1	0	9
2	9960.0	3	2	2	2	0	1	1	0	6
3	7500.0	4	2	2	3	0	1	1	0	7
4	7420.0	4	1	2	2	0	1	0	1	7

```
In [66]: # dependent variables
y = df[['price']]
y.head()
```

```
Out[66]:
```

	price
0	4340000.0
1	4340000.0
2	4340000.0
3	4340000.0
4	11410000.0

9. Scaling the independent variables

```
In [67]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_std = scaler.fit_transform(X)
```

```
In [68]: x_std
```

```
Out[68]: array([[ 1.11756482,  1.40341936,  1.42181174, ...,  1.70084013,
        -0.84488844, -0.6964292 ],
        [ 1.8623093 ,  1.40341936,  5.40580863, ...,  1.70084013,
        -0.84488844, -0.6964292 ],
        [ 2.34590961,  0.04727831,  1.42181174, ..., -0.58794474,
         1.18358821, -0.6964292 ],
        ...,
        [-0.72011635, -1.30886273, -0.57018671, ..., -0.58794474,
        -0.84488844,  1.43589615],
        [-1.06347257,  0.04727831, -0.57018671, ...,  1.70084013,
        -0.84488844, -0.6964292 ],
        [-0.60888828,  0.04727831, -0.57018671, ..., -0.58794474,
        -0.84488844,  1.43589615]])
```

10. Split the data into training and testing

```
In [69]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=0)
```

11. Build the Model

```
In [70]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [71]: lr = LinearRegression()
```

```
In [ ]:
```

12. Train the Model

```
In [72]: lr.fit(X_train, y_train)
```

```
Out[72]: ▼ LinearRegression
LinearRegression()
```

```
In [73]: print("Value of the coefficients: \n", lr.coef_)
print(" ----- ")
print("Value of the intercept: \n", lr.intercept_)
```

```
Value of the coefficients:
[[ 2.60781675e+02  9.30932038e+04  8.20605321e+05  3.96961106e+05
  1.00796216e+05 -3.10469714e+05  3.10469714e+05 -3.32132424e+05
  3.32132424e+05 -1.72635846e+05  1.72635846e+05 -7.33905647e+05
  7.33905647e+05 -4.83024979e+05  4.83024979e+05  1.04827468e+05
  1.38634062e+05 -2.43461530e+05]]
-----
Value of the intercept:
[2156139.9017023]
```

13. Test the Model

```
In [74]: Y_pred = lr.predict(X_test)
```

14. Measure the performance using Metrics.

```
In [81]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
In [82]: print("MSE",mean_squared_error(y_test,Y_pred))
print(" ----- ")
print("RMSE",np.sqrt(mean_squared_error(y_test,Y_pred)))
print(" ----- ")
print("R-Square", r2_score(y_test,Y_pred))
```

```
MSE 1326521791171.129
-----
RMSE 1151747.2774750236
-----
R-Square 0.5718914765881087
```

```
In [ ]:
```